

Fundamentals of Deep Learning and Programming

Lecture 4: Neural Network Training


Fall Semester 2026

Adjunct Professor: Donghoon Kang

Learn from Data



Future Data



Train



Past Data
(perhaps with **labels**)

Supervised, Semi-Supervised, and Unsupervised Learning

Training data

All data is **labeled**

Supervised
Learning

**Learn a function f to
map data x to label y :**
 $y = f(x)$

Small portion of
data is **labeled**

Lots of data is
unlabeled

Semi-Supervised
Learning

All data is **unlabeled**

Unsupervised
Learning

**Learn some underlying
hidden structure of
data**

Problem: Supervised Learning is **NOT** how we humans learn



Baby **don't get supervision** for everything they see !

Self-Supervised Learning

- Self-supervised learning is a machine learning paradigm where the model learns from unlabeled data by creating its own supervised signals.
 - In self-supervised learning, the model generates “labels” automatically from the raw data and then trains itself as if it were a supervised task.
 - Self-supervised learning is generally considered unsupervised learning in terms of **data** (no manual labels), but supervised learning in terms of **method** (training with input–output prediction tasks).

Rapid Generalization and Efficient Learning like Human

Traditionally, a deep learning model needs thousands, or even millions, of examples to learn a new task. But what if we don't have enough data?

- This is where Zero-shot, One-shot, and Few-shot learning come in. These concepts are inspired by how humans learn—we can often recognize a new concept with very few, or sometimes zero, examples.

Zero-shot, One-shot, and Few-shot leverage **transfer learning** and **meta-learning (=learning to learn)** to enable models to generalize with minimal data, a major step towards more intelligent AI.

Zero-Shot Learning (ZSL)


The model should classify an object from a class it has never seen before during training.

One-Shot Learning (OSL)

The model is given only one example of a new class and must classify a new instance of that class.

Few-Shot Learning (FSL)


The model is given a small number (e.g., 3-5) of examples per new class and must classify a new instance of that class.



Zero-Shot Learning

Zero-shot learning means the model can recognize a new class **without ever seeing any training example** of it. Instead, it **uses semantic descriptions or prior knowledge**.

Example: The model has never seen a zebra photo. But if it knows that a zebra is “a horse-like animal with black and white stripes”, it can use this description to classify zebras correctly, even without any zebra training image.




Another Examples:

- Machine Translation: Translating from English to German without having seen parallel training data for that specific language pair.
- Prompting LLMs: Asking ChatGPT to solve a task without examples in the prompt.

One-Shot Learning


One-shot learning means the model can correctly recognize a new class after seeing **just one example** of it.

Example: Suppose the model has never seen a zebra before. If we show the model **a single** photo of a zebra, it should learn to **generalize** and correctly **classify** other zebra images.



Another **Example:**


Face recognition: Given **one picture** of a person, the model can recognize them in other images.



Few-Shot Learning


In Few-shot learning, the model is given a few examples (e.g., 5–10 images) of the new class before making predictions.

Example: If we provide the model with 3 photos of zebras from different angles, it should then be able to identify zebras in new, unseen images.




Additional Examples:

- Medical imaging: Classifying rare diseases where only a handful of labeled cases exist.



- LLMs with prompts: If you give GPT a few examples in the prompt (“few-shot prompting”), it improves accuracy on the task.

Perceptron



1. Introduction

The **Perceptron** is one of the earliest models of an artificial neuron, introduced by Frank Rosenblatt in 1958. It is a simple binary linear classifier that inspired the development of modern neural networks.

2. Model Definition

Given an input vector

$$\mathbf{x} = (x_1, x_2, \dots, x_n)^T \in \mathbb{R}^n,$$

the perceptron computes a weighted sum with bias:

$$z = \mathbf{w}^T \mathbf{x} + b,$$

where $\mathbf{w} = (w_1, w_2, \dots, w_n)^T$ is the weight vector and b is the bias term.

The perceptron output is obtained by applying a step activation function:

$$y = f(z) = \begin{cases} 1, & z \geq 0, \\ 0, & z < 0. \end{cases}$$

Perceptron

3. Learning Rule

The perceptron updates its weights iteratively to reduce classification errors. For a training sample $(\mathbf{x}, t)$ with target label $t \in \{0, 1\}$:

$$\mathbf{w} := \mathbf{w} + \eta(t - y)\mathbf{x},$$

$$b := b + \eta(t - y),$$


where $\eta > 0$ is the learning rate.

4. Properties


- The perceptron can classify linearly separable data.
- If the data is linearly separable, the perceptron learning algorithm converges in a finite number of steps.
- It cannot solve problems that are not linearly separable (e.g., the XOR problem).
- Despite its limitations, the perceptron was the foundation for modern neural network models.

5. Extension to Neural Networks

A single perceptron corresponds to one neuron. By stacking multiple perceptrons in layers and using differentiable activation functions (e.g., sigmoid, ReLU), we obtain modern multilayer neural networks capable of solving non-linear problems.



Neural Network (e.g., classification)



Hidden Layer

For each hidden neuron h_j ($j = 1, 2, 3, 4$):

$$z_j^{[1]} = w_{j1}^{[1]}x_1 + w_{j2}^{[1]}x_2 + w_{j3}^{[1]}x_3 + b_j^{[1]}$$

$$h_j = f(z_j^{[1]})$$

where $f(\cdot)$ is an activation function (e.g., ReLU, sigmoid).

Activation Function (1): Sigmoid Function

```
import numpy as np
x = np.linspace(-5, 5, 100)
sigmoid = 1 / (1 + np.exp(-x))
```

Numpy


```
import torch
x = torch.linspace(-5, 5, 100)
sigmoid = torch.sigmoid(x)
```

PyTorch

$$f(x) = \frac{1}{1 + e^{-x}}$$

$$\lim_{x \rightarrow -\infty} f(x) = 0, \quad \lim_{x \rightarrow +\infty} f(x) = 1$$

- Maps any real number into the range $[0, 1]$.
- Smooth and differentiable everywhere.
- Saturates for very large positive or negative inputs, which can cause the vanishing gradient problem.




Activation Function (2): Tanh (Hyperbolic Tangent)

```
import numpy as np
x = np.linspace(-5, 5, 100)
tanh = np.tanh(x)
```

Numpy

```
import torch
x = torch.linspace(-5, 5, 100)
tanh = torch.tanh(x)
```

PyTorch



$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\lim_{x \rightarrow -\infty} f(x) = -1, \quad \lim_{x \rightarrow +\infty} f(x) = 1$$

- Maps any real number into the range $[-1, 1]$.
- Zero-centered, unlike the sigmoid function.
- Still suffers from vanishing gradients for very large positive or negative inputs.

Activation Function (3): ReLU (Rectified Linear Unit)

```
import numpy as np
x = np.linspace(-5, 5, 100)
relu = np.maximum(0, x)
```


Numpy

```
import torch
x = torch.linspace(-5, 5, 100)
relu = torch.relu(x)
```

PyTorch

$$f(x) = \max(0, x)$$

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ 0, & \text{if } x \leq 0 \end{cases}$$



- Simple and computationally efficient.
- Helps avoid the vanishing gradient problem compared to sigmoid or tanh.
- Not differentiable at $x = 0$, but works well in practice.

Activation Function (4): GELU (Gaussian Error Linear Unit)

1) Definition


$$f(x) = x \cdot \Phi(x)$$

where $\Phi(x)$ is the cumulative distribution function (CDF) of the standard Gaussian distribution:

$$\Phi(x) = \frac{1}{2} \left(1 + \operatorname{erf} \left(\frac{x}{\sqrt{2}} \right) \right)$$

Thus, explicitly:

$$f(x) = \frac{x}{2} \left(1 + \operatorname{erf} \left(\frac{x}{\sqrt{2}} \right) \right)$$




2) Approximation

Since computing the error function can be expensive, the following tanh-based approximation is often used:

$$f(x) \approx 0.5x \left(1 + \tanh \left(\sqrt{\frac{2}{\pi}} (x + 0.044715x^3) \right) \right)$$

Activation Function (4): GELU (Gaussian Error Linear Unit)



- Smooth and differentiable, unlike ReLU which has a sharp cutoff at zero.
- Probabilistically scales negative inputs instead of discarding them completely.
- Widely used in modern deep learning architectures such as BERT, GPT, and Vision Transformers.

Activation Function (4): GELU (Gaussian Error Linear Unit)

```
import numpy as np
from math import sqrt, pi
from math import erf # scalar version
# For vectorized erf, scipy.special.erf is recommended

def gelu_erf_numpy(x: np.ndarray) -> np.ndarray:
    # Exact GeLU:  $f(x) = 0.5 * x * (1 + \text{erf}(x / \sqrt{2}))$ 
    return 0.5 * x * (1.0 + np.vectorize(erf)(x / sqrt(2.0)))

def gelu_tanh_numpy(x: np.ndarray) -> np.ndarray:
    # Approximate GeLU:
    #  $f(x) \approx 0.5 * x * (1 + \tanh(\sqrt{2/\pi}) * (x + 0.044715 * x^3))$ 
    return 0.5 * x * (1.0 + np.tanh(np.sqrt(2.0 / np.pi) * (x + 0.044715 * x**3)))

# Example usage
x = np.linspace(-5, 5, 100)
y_exact = gelu_erf_numpy(x)
y_approx = gelu_tanh_numpy(x)
```

Numpy

Activation Function (4): GELU (Gaussian Error Linear Unit)

```
import torch
import torch.nn as nn
import torch.nn.functional as F


# Functional API
y_exact = F.gelu(x, approximate='none')      # Exact (ERF-based) version
y_approx = F.gelu(x, approximate='tanh')      # Tanh-based approximation

# Module API
gelu_exact = nn.GELU(approximate='none')
gelu_approx = nn.GELU(approximate='tanh')

y1 = gelu_exact(x)      # y_exact == gelu_exact
y2 = gelu_approx(x)     # y_approx == gelu_approx
```

PyTorch

Neural Network (e.g., classification)




Output Layer

For each output neuron $\hat{y}_k$ ($k = 1, 2$):

$$z_k^{[2]} = w_{k1}^{[2]}h_1 + w_{k2}^{[2]}h_2 + w_{k3}^{[2]}h_3 + w_{k4}^{[2]}h_4 + b_k^{[2]}$$

$$\hat{y}_k = \frac{e^{z_k^{[2]}}}{\sum_{l=1}^2 e^{z_l^{[2]}}} \quad (\text{Softmax activation})$$

Neural Network (e.g., classification)




Compact Matrix Form

$$\mathbf{h} = f(W^{[1]}\mathbf{x} + \mathbf{b}^{[1]}), \quad \mathbf{x} \in \mathbb{R}^3, \mathbf{h} \in \mathbb{R}^4$$

$$\hat{\mathbf{y}} = \text{Softmax}(W^{[2]}\mathbf{h} + \mathbf{b}^{[2]}), \quad \hat{\mathbf{y}} \in \mathbb{R}^2$$

Neural Network (e.g., classification)



Loss Function (Cross-Entropy)

For target label $y \in \{1, 2\}$:

$$L = - \sum_{k=1}^2 \mathbf{1}\{y = k\} \log(\hat{y}_k)$$

where $\mathbf{1}\{y = k\}$ is an indicator function equal to 1 if the true class is k , otherwise 0.

Softmax Function

The **softmax function** converts a vector of real numbers into a probability distribution. Given an input vector

$$\mathbf{z} = (z_1, z_2, \dots, z_K) \in \mathbb{R}^K,$$

the softmax function is defined as:

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}, \quad \text{for } i = 1, 2, \dots, K.$$

- Each output value lies in the range $(0, 1)$.
- The outputs sum to 1:

$$\sum_{i=1}^K \sigma(\mathbf{z})_i = 1.$$

- Therefore, the softmax function transforms raw scores (logits) into a probability distribution over K classes.
- Softmax is commonly used in the final layer of a neural network for multi-class classification tasks.

Softmax Function

```
import torch

logits = torch.tensor([[1.0, 2.0, 3.0],
                       [1.0, 2.0, 5.0]])

a = torch.softmax(logits, dim=0) # column-wise
b = torch.softmax(logits, dim=1) # row-wise

print(a)
print(a.sum(dim=0)) # should be all ones

print(b)
print(b.sum(dim=1)) # should be all ones
```

```
tensor([[0.5000, 0.5000, 0.1192],
        [0.5000, 0.5000, 0.8808]])
tensor([1.0000, 1.0000, 1.0000])
```

```
tensor([[0.0900, 0.2447, 0.6652],
        [0.0171, 0.0466, 0.9362]])
tensor([1., 1.])
```

```
a = torch.softmax(logits, dim=0)
→ normalizes down the rows, so each column sums to 1.

b = torch.softmax(logits, dim=1)
→ normalizes across the columns, so each row sums to 1.
```

Sigmoid Function in the output layer

1. Sigmoid Function

The sigmoid function is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

It maps any real number into the range $(0, 1)$, which makes it suitable for binary classification.

- Binary classification ($K = 2$) $\rightarrow$ use **sigmoid**.
- Multi-class classification ($K > 2$) $\rightarrow$ use **softmax**.

2. Binary Classification (2 Classes)

For binary classification, the neural network produces a single logit z . Applying the sigmoid gives:

$$\hat{y} = \sigma(z) \in (0, 1)$$

Interpretation:

$$P(\text{Class 1} \mid z) = \hat{y}, \quad P(\text{Class 0} \mid z) = 1 - \hat{y}$$

The loss function is the **Binary Cross-Entropy (BCE)**:

$$L = -\left(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})\right), \quad y \in \{0, 1\}$$

torch.nn.Linear ()

```
import torch
import torch.nn as nn

# Create a linear layer:  $y = xW + b$ 
layer = nn.Linear(in_features=4, out_features=2)

# Random input tensor with shape (batch_size=3, in_features=4)
x = torch.randn(3, 4)

# Call the layer like a function
y = layer(x)

print("Input shape:", x.shape)
print("Output shape:", y.shape)
```

```
Input shape: torch.Size([3, 4])
Output shape: torch.Size([3, 2])
```

`torch.nn.Linear` are **callable** in PyTorch.

This is because **all subclasses of `torch.nn.Module` implement** the special method **`__call__`**, which internally calls the forward method.

- `nn.Linear` is a class → you must instantiate it (e.g., `layer = nn.Linear(4, 2)`).
- The instance (`layer`) is callable like a function because it inherits from `nn.Module`.
- `layer.__call__()` → does pre/post processing → calls `layer.forward()`.

1. Purpose of Initialization

When training deep neural networks:

- If weights are too **large**, activations and gradients may **explode**.
- If weights are too **small**, activations and gradients may **vanish**.

Good initialization tries to keep the **variance of activations and gradients stable** across layers.

2. Target Variance

We define a **target variance** v for the weights:

$$\text{Var}(W) = v.$$

The value of v depends on the chosen initialization scheme:

- For **Kaiming (He)** initialization (ReLU),

$$v = \frac{2}{\text{fan_in}}.$$

- For **Xavier (Glorot)** initialization,

$$v = \frac{2}{\text{fan_in} + \text{fan_out}}.$$

All parameter derivations (uniform bound, Gaussian variance) aim to enforce this condition.

3. Why a Bound or Variance is Used

There are two common sampling strategies:

- **Uniform distribution:**

$$W \sim \mathcal{U}(-\text{bound}, +\text{bound}),$$

where the *bound* is chosen so that $\text{Var}(W) = v$.

- **Normal (Gaussian) distribution:**

$$W \sim \mathcal{N}(0, \sigma^2),$$

where the variance σ^2 is directly set to v .

4. Example: Variance of Uniform vs Normal

Uniform distribution. For $X \sim U(-a, a)$:

$$\text{Var}(X) = \frac{a^2}{3}.$$

Setting $\text{Var}(X) = v$ gives

$$a = \sqrt{3v}.$$

Normal distribution. For $X \sim \mathcal{N}(0, \sigma^2)$:

$$\text{Var}(X) = \sigma^2.$$

So we directly set

$$\sigma^2 = v.$$

5. In Kaiming and Xavier Initialization (Weights)

- **Kaiming (He) initialization (ReLU case):**

$$v = \frac{2}{\text{fan_in}}.$$

- Uniform form:

$$\text{bound} = \sqrt{\frac{6}{\text{fan_in}}}.$$

- Normal form:

$$\sigma^2 = \frac{2}{\text{fan_in}}.$$

- **Xavier (Glorot) initialization:**

$$v = \frac{2}{\text{fan_in} + \text{fan_out}}.$$

- Uniform form:

$$\text{bound} = \sqrt{\frac{6}{\text{fan_in} + \text{fan_out}}}.$$

- Normal form:

$$\sigma^2 = \frac{2}{\text{fan_in} + \text{fan_out}}.$$

6. Bias Initialization in PyTorch

In contrast to weights, PyTorch initializes biases using a simpler rule:

$$b \sim \mathcal{U}\left(-\frac{1}{\sqrt{\text{fan_in}}}, +\frac{1}{\sqrt{\text{fan_in}}}\right).$$

This does not aim to preserve variance, but simply keeps bias values small relative to the number of input connections.

- **Kaiming (He) Initialization**

- Default for `nn.Linear` in PyTorch.
- Uses only `fan_in`.
- Best for ReLU-like activations.

- **Xavier (Glorot) Initialization**

- Must be applied manually with `torch.nn.init`.
- Uses both `fan_in` and `fan_out`.
- Best for sigmoid/tanh activations.

7. PyTorch Example (Conceptual)

```
import torch, math

in_features, out_features = 3, 2
fan_in, fan_out = in_features, out_features

# Target variance v
v_kaiming = 2.0 / fan_in
v_xavier = 2.0 / (fan_in + fan_out)

# Kaiming initialization
bound_kaiming = math.sqrt(3 * v_kaiming)           # uniform
sigma_kaiming = math.sqrt(v_kaiming)              # normal

# Xavier initialization
bound_xavier = math.sqrt(3 * v_xavier)            # uniform
sigma_xavier = math.sqrt(v_xavier)               # normal

# Bias bound
bound_bias = 1.0 / math.sqrt(fan_in)

# Generate weights and bias
w_kaiming_u = torch.empty(out_features, in_features).uniform_(-bound_kaiming, bound_kaiming)
w_kaiming_n = torch.empty(out_features, in_features).normal_(0.0, sigma_kaiming)

w_xavier_u = torch.empty(out_features, in_features).uniform_(-bound_xavier, bound_xavier)
w_xavier_n = torch.empty(out_features, in_features).normal_(0.0, sigma_xavier)

b_init = torch.empty(out_features).uniform_(-bound_bias, bound_bias)
```

```
import torch
import torch.nn as nn

layer = nn.Linear(3, 2)

print(layer.weight)
print(layer.bias)
```

Check Kaiming and Xavier Initialization

```
import torch
import torch.nn as nn
import torch.nn.init as init
import math

# Define a linear layer
layer = nn.Linear(in_features=3, out_features=2, bias=True)

# --- Default behavior ---
# By default, nn.Linear uses Kaiming uniform initialization for weights
# and uniform(-1/sqrt(fan_in), 1/sqrt(fan_in)) for bias.
print("Default (Kaiming) weight:\n", layer.weight)
print("Default bias:\n", layer.bias)

# --- Xavier initialization (optional) ---
# If you prefer Xavier, reinitialize explicitly:
init.xavier_uniform_(layer.weight)
# The trailing underscore _ in PyTorch convention means
# in-place operation (it modifies the tensor directly)

fan_in = layer.in_features
bound_bias = 1 / math.sqrt(fan_in)

if layer.bias is not None:
    layer.bias.data.uniform_(-bound_bias, bound_bias)

print("\nXavier-initialized weight:\n", layer.weight)
print("Xavier-initialized bias:\n", layer.bias)
```

```
Default (Kaiming) weight:
Parameter containing:
tensor([[ -0.0480,  0.2426, -0.2729],
        [ 0.5431,  0.3103,  0.0399]], requires_grad=True)
Default bias:
Parameter containing:
tensor([0.3220, 0.4163], requires_grad=True)

Xavier-initialized weight:
Parameter containing:
tensor([[ 0.7014,  0.7367,  0.1130],
        [-0.4929, -0.0240, -0.5067]], requires_grad=True)
Xavier-initialized bias:
Parameter containing:
tensor([ 0.0247, -0.2385], requires_grad=True)
```

Backpropagation


Linear Regression

Input layer
(2 nodes)

hidden layer
(3 nodes)

Output layer
(2 nodes)

Activation Functions: **ReLU**



Backpropagation – Linear Regression (2-3-2)

```
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
```

```
# ---- 1) Model ----
```


```
class myModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(2, 3, bias=True) # input(2) -> hidden(3)
        self.fc2 = nn.Linear(3, 2, bias=True) # hidden(3) -> output(2)
        self.reset_to_given_values()
```

```
    def reset_to_given_values(self):
        # Given matrices in the document
        W_ih = torch.tensor([[0.1, 0.2, 0.3],
                              [0.4, 0.5, 0.6]], dtype=torch.float32)
        B_h = torch.tensor([0.2, 0.4, 0.6], dtype=torch.float32)
        W_ho = torch.tensor([[0.7, 0.8],
                              [0.9, 0.1],
                              [0.2, 0.3]], dtype=torch.float32)
        B_o = torch.tensor([0.5, 0.7], dtype=torch.float32)
```

```
    with torch.no_grad():
        self.fc1.weight.copy_(W_ih.t()) # (3,2)
        self.fc1.bias.copy_(B_h)
        self.fc2.weight.copy_(W_ho.t()) # (2,3)
        self.fc2.bias.copy_(B_o)
```

```
    def forward(self, x):
        x = torch.relu(self.fc1(x)) # hidden layer only
        x = self.fc2(x) # output layer: no ReLU
        return x
```

c03_backprop_linear_reg_ReLU_242.py



IMPORTANT




Remove the ReLU from the output layer

- When you used the **default PyTorch initialization**, it's possible that the **pre-activation outputs** of the final layer were **negative**.
- If they are negative, the output layer's ReLU sets them to zero, and the gradient of ReLU for negative inputs is also zero.
- Result: those output units become "dead" → no gradient flows → weights don't update → the loss does not decrease.

$$f(x) = \max(0, x)$$

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ 0, & \text{if } x \leq 0 \end{cases}$$




Remove the ReLU from the output layer

- For regression tasks, it's still best practice to **remove the ReLU from the output layer**.
- But if your initialization (i.e., lucky case) happens to place outputs in the ReLU's active region, training can still succeed.

$$f(x) = \max(0, x)$$

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ 0, & \text{if } x \leq 0 \end{cases}$$



How forward() is called in PyTorch

- In PyTorch, you don't usually call forward() directly.

```
import torch
import torch.nn as nn

class MyModel(nn.Module):
    def forward(self, x):
        print('hello')
        return x

x = torch.tensor([1.0])

net = MyModel()           # create an instance
y = net(x)                # calls forward(x)
print(y)
```


- MyModel() triggers the `__call__` method defined in the base class **nn.Module**.
- Inside `__call__`, PyTorch does several things automatically:
 - ✓ Prepares the input for autograd (automatic differentiation).
 - ✓ Calls your model's forward() method with the input.
 - ✓ Processes the output and attaches the computation graph for backpropagation.
- So, in effect, model(x) → calls forward(x) behind the scenes.

Backpropagation – Linear Regression (2-3-2)

```
# ---- 2) Data / Loss / Optimizer ----
model = myModel()
x = torch.tensor([0.5, 0.8], dtype=torch.float32)
t = torch.tensor([0.1, 0.9], dtype=torch.float32)

criterion = nn.MSELoss(reduction='mean') # default MSE
learning_rate = 0.10
optimizer = optim.SGD(model.parameters(), lr=learning_rate)

# ---- 3) print parameters ----
def print_params(step, mdl):
    with torch.no_grad():
        Wih = mdl.fc1.weight.t()
        Bh = mdl.fc1.bias
        Who = mdl.fc2.weight.t()
        Bo = mdl.fc2.bias
        print(f"\n[After step {step}]")
        print("W_ih =\n", Wih.numpy())
        print("B_h = ", Bh.numpy())
        print("W_ho =\n", Who.numpy())
        print("B_o = ", Bo.numpy())
```



Backpropagation – Linear Regression (2-3-2)


```
# ---- 4) Training loop with loss tracking ----
losses = []

for step in range(1, 21): # 20 steps to see the curve
    optimizer.zero_grad()
    y = model(x)
    loss = criterion(y, t)
    loss.backward()
    optimizer.step()

    losses.append(loss.item()) # record loss
    print(f"step {step} | loss = {loss.item():.6f}")

    if step % 5 == 0:
        print_params(step, model)

# ---- 5) Plot loss curve ----
plt.figure(figsize=(6,4))
plt.plot(range(1, len(losses)+1), losses, marker='o')
plt.xlabel("Step")
plt.ylabel("Loss (MSE)")
plt.title("Loss Curve during Training")
plt.grid(True)
plt.show()
```



```
[After step 5]
W_ih =
[[ -0.04336868  0.05882919  0.27575284]
 [ 0.17061011  0.2741267   0.5612046 ]]
B_h = [-0.08673736  0.11765838  0.5515057 ]
W_ho =
[[ 0.571718    0.7558604 ]
 [ 0.6565286   0.02737943]
 [-0.19934106  0.19679347]]
B_o = [0.16177835 0.6148201 ]
step 6 | loss = 0.004628
step 7 | loss = 0.002013
step 8 | loss = 0.000912
step 9 | loss = 0.000425
step 10 | loss = 0.000202
```

Compact code for Linear Regression (2-3-2)


```
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

# ---- 1) Model ----
class MyModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(2, 3) # input(2) -> hidden(3)
        self.fc2 = nn.Linear(3, 2) # hidden(3) -> output(2)

    def forward(self, x):
        x = torch.relu(self.fc1(x)) # ReLU only in the hidden layer
        x = self.fc2(x)             # No activation on the output layer
        return x

# ---- 2) Data / Loss / Optimizer ----
model = MyModel()
x = torch.tensor([0.5, 0.8], dtype=torch.float32).unsqueeze(0) # [1, 2]
t = torch.tensor([0.1, 0.9], dtype=torch.float32).unsqueeze(0) # [1, 2]

criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01) # smaller LR for stability
```



C03_linear_reg_ReLU_242.py



Compact code for Linear Regression (2-3-2)

```
# ---- 3) Training loop with loss tracking ----
losses = []
for step in range(1, 101): # 100 steps for clearer trend
    optimizer.zero_grad()
    y = model(x)
    loss = criterion(y, t)
    loss.backward()
    optimizer.step()


    losses.append(loss.item())
    if step % 10 == 0:
        print(f"step {step:3d} | loss = {loss.item():.6f}")

# ---- 4) Plot loss curve ----
plt.figure(figsize=(6,4))
plt.plot(range(1, len(losses)+1), losses, marker='o')
plt.xlabel("Step")
plt.ylabel("Loss (MSE)")
plt.title("Loss Curve during Training")
plt.grid(True)
plt.show()
```

step 10	loss = 0.403117
step 20	loss = 0.329712
step 30	loss = 0.269674
step 40	loss = 0.220568
step 50	loss = 0.180404
step 60	loss = 0.147554
step 70	loss = 0.120685
step 80	loss = 0.098709
step 90	loss = 0.080735
step 100	loss = 0.066034



Activation Functions: **ReLU**



Backpropagation computes gradients efficiently using the chain rule. We use **ReLU** in the hidden layer and **Softmax + Cross-Entropy** at the output for a 4-class classification task.

$$\text{ReLU}(z) = \max(0, z), \quad \text{ReLU}'(z) = \begin{cases} 1, & z > 0, \\ 0, & z \leq 0. \end{cases}$$

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}, \quad E = - \sum_{i=1}^4 t_i \log(o_i).$$

We keep the order *local gradient* $\times$ *upstream gradient* throughout, and denote error signals by q_o (output) and q_h (hidden).

```
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
```

Backpropagation

```
torch.set_printoptions(precision=6, sci_mode=False)
```

```
# ---- 1) Model ----
```


```
class MyModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(2, 3, bias=True) # input(2) -> hidden(3)
        self.fc2 = nn.Linear(3, 4, bias=True) # hidden(3) -> output(4) (logits)
        self.reset_to_given_values()

    def reset_to_given_values(self):
        W_ih = torch.tensor([[0.1, 0.2, 0.3],
                               [0.4, 0.5, 0.6]], dtype=torch.float32) # (2x3)
        B_h = torch.tensor([0.2, 0.4, 0.6], dtype=torch.float32) # (3,)
        W_ho = torch.tensor([[0.7, 0.8, 0.5, 0.2],
                               [0.9, 0.1, 0.4, 0.3],
                               [0.2, 0.3, 0.6, 0.7]], dtype=torch.float32) # (3x4)
        B_o = torch.tensor([0.5, 0.7, 0.2, 0.1], dtype=torch.float32) # (4,)

        with torch.no_grad():
            # nn.Linear stores weight as (out_features, in_features)
            self.fc1.weight.copy_(W_ih.t()) # (3,2)
            self.fc1.bias.copy_(B_h)
            self.fc2.weight.copy_(W_ho.t()) # (4,3)
            self.fc2.bias.copy_(B_o)

    def forward(self, x): # x: (N,2)
        h = torch.relu(self.fc1(x))
        logits = self.fc2(h) # raw scores for CrossEntropyLoss
        return logits
```

Classification



Backpropagation

Classification

```
# ---- 2) Data / Loss / Optimizer ----
```

```
model = MyModel()
```

```
# Batch of one sample to play nicely with CrossEntropyLoss
```

```
x = torch.tensor([[0.5, 0.8]], dtype=torch.float32) # (1,2)
```

```
target_index = torch.tensor([3], dtype=torch.long) # (1,)
```

```
criterion = nn.CrossEntropyLoss(reduction='mean') # applies log-softmax + NLL
```

```
learning_rate = 0.10
```

```
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
```

```
# ---- 3) print parameters ----
```

```
def print_params(step, mdl):
```

```
    with torch.no_grad():
```

```
        Wih = mdl.fc1.weight.t()
```

```
        Bh = mdl.fc1.bias
```

```
        Who = mdl.fc2.weight.t()
```

```
        Bo = mdl.fc2.bias
```

```
        print(f"\n[After step {step}]\n")
```

```
        print("W_ih =\n", Wih.numpy())
```

```
        print("B_h = ", Bh.numpy())
```


```
        print("W_ho =\n", Who.numpy())
```

```
        print("B_o = ", Bo.numpy())
```

```
# Helper to view softmax probs
```

```
def softmax_probs(logits):
```

```
    return torch.softmax(logits, dim=1) # softmax per row (across classes) --> rows sum to 1
```



```
# ---- 4) Training loop with loss tracking ----
```

```
losses = []
```

```
for step in range(1, 21): # 20 steps to see the curve
    optimizer.zero_grad()
    logits = model(x) # (1,4)
    probs = softmax_probs(logits) # for inspection only
    loss = criterion(logits, target_index) # CE on logits
    loss.backward()
```

```
# (Optional) On step 1, verify q_o = probs - one_hot equals grad wrt output bias
```

```
if step == 1:
    with torch.no_grad():
        one_hot = torch.zeros_like(probs)
        one_hot[0, target_index.item()] = 1.0
        q_o = probs - one_hot # shape (1,4)
        print("\n[Step 1: gradient check at output layer]")
        print("logits =", logits.detach().squeeze(0))
        print("probs  =", probs.detach().squeeze(0))
        print("q_o (= probs - one_hot)          =", q_o.squeeze(0))
        print("grad wrt output bias (fc2.bias)=", model.fc2.bias.grad)
```


```
optimizer.step()
```

```
losses.append(loss.item())
print(f"step {step} | loss = {loss.item():.6f} | probs = {probs.detach().numpy().round(6)}")
```

```
if step % 5 == 0:
    print_params(step, model)
```

Backpropagation

Classification





Backpropagation

Classification

```
# ---- 5) Plot loss curve ----  
plt.figure(figsize=(6,4))  
plt.plot(range(1, len(losses)+1), losses, marker='o')  
plt.xlabel("Step")  
plt.ylabel("Loss (CrossEntropy)")  
plt.title("Loss Curve during Training (Softmax + CE)")  
plt.grid(True)  
plt.show()
```

```
[After step 5]  
W_ih =  
[[-0.04336868  0.05882919  0.27575284]  
 [ 0.17061011  0.2741267   0.5612046  ]]  
B_h = [-0.08673736  0.11765838  0.5515057 ]  
W_ho =  
[[ 0.571718    0.7558604 ]  
 [ 0.6565286   0.02737943]  
 [-0.19934106  0.19679347]]  
B_o = [0.16177835  0.6148201 ]  
step 6 | loss = 0.004628  
step 7 | loss = 0.002013  
step 8 | loss = 0.000912  
step 9 | loss = 0.000425  
step 10 | loss = 0.000202
```



Thank you!

